

Міністерство освіти і науки України
Національний технічний університет України
“Київський політехнічний інститут”
Кафедра АСОІУ

ЗВІТ
про виконання лабораторної роботи №6
з дисципліни
“ Аналіз даних в інформаційно-управляючих системах”

Виконав Студент
2 курсу групи ІП-11
Панченко Сергій

Київ 2023

Lab Six

Main Assignment

```
In [39]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import plotly.express as px
import missingno as msno
from kneed import KneeLocator
from sklearn.cluster import KMeans, AgglomerativeClustering
import math
```

Correct File

```
In [40]: with open('data/titanic.csv') as input_file:
    with open('data/data_titanic.csv', 'w') as output_file:
        lines = []
        statuses = []
        for index, line in enumerate(input_file):
            if index == 0:
                lines.append(line.replace(',', ';'))
                continue
            parts = line.split('"')
            parts[0] = parts[0].replace(',', ';')
            parts[2] = parts[2].replace(',', ';')
            statuses.append(parts[1].split(',')[1].split('.')[0].lstrip())
            lines.append(''.join(parts))
        statuses = list(set(statuses))
        output_file.writelines(lines)
```

Status-Gender Pairs

```
In [41]: status_gender_pairs = {
    'Rev': 'male',
    'Miss': 'female',
    'Ms': 'male',
    'Dr': 'male',
    'Mlle': 'female',
    'Col': 'male',
    'the Countess': 'female',
    'Jonkheer': 'male',
    'Don': 'male',
    'Capt': 'male',
    'Master': 'male',
    'Mme': 'female',
    'Mrs': 'female',
    'Mr': 'male',
    'Lady': 'female',
    'Sir': 'male',
    'Major': 'male'
}
```

Load Dataframe

```
In [42]: df = pd.read_csv('data/data_titanic.csv', delimiter=';', decimal='.')
```

Drop unnecessary columns

```
In [43]: df = df.drop(['PassengerId'], axis=1)
```


Take a look on what data is missing

```
In [44]: percent_1 = df.isnull().sum() / df.isnull().count() * 100
percent_2 = (round(percent_1, 1)).sort_values(ascending=False)
missing_data = pd.concat([percent_2], axis=1, keys=['%'])
plot_rows = 1
plot_cols = 1
fig, axis = plt.subplots(plot_rows, plot_cols, figsize=(8, 6))
msno.matrix(df, ax=axis)
```

/usr/local/lib/python3.10/dist-packages/missingno/missingno.py:61: UserWarning:

Plotting a sparkline on an existing axis is not currently supported. To remove this warning, set sparkline=False.


Out[44]: <AxesSubplot: >



Empty values by percentage in column

In [45]: fig, axis = plt.subplots(plot_rows, plot_cols, figsize=(8, 6))
axis.bar(missing_data.index, missing_data.iloc[:, 0])

Out[45]: <BarContainer object of 11 artists>



Drop "Cabin" due to large number of missing data

```
In [46]: df = df.drop(['Cabin'], axis=1)
```

Fill Nan "Age"-values with mean

```
In [47]: df_sex_mean_age = pd.pivot_table(df, index='Sex', values='Age', aggfunc='mean')
mean_age_female, mean_age_male = df_sex_mean_age.iloc[0, 0], df_sex_mean_age.iloc[1, 0]
df_sex_mean_age
```

```
Out[47]:
```

	Age
Sex	
female	28.445833
male	30.863713

Fill gender

```
In [48]: df['Sex'] = df.apply(lambda x: status_gender_pairs[x['Name'].split(',')[1].split('.')[0].lstrip()], axis=1)
df['Gender'] = df.apply(lambda x: x['Sex'] == 'male', axis=1)


def mean_age_func(row):
    if str(row['Age']) != 'nan':
        return row['Age']
    if row['Sex'] == 'male':
        return mean_age_male
    return mean_age_female

df['Age'] = df.apply(mean_age_func, axis=1)
```

Fill "Embarked" with the most frequent value

```
In [49]: fig, axis = plt.subplots(figsize=(8, 6))
embarked_counts = df['Embarked'].value_counts()
axis.pie(embarked_counts, labels=embarked_counts.index, autopct='%1.1f%%')
```

```
Out[49]: ([<matplotlib.patches.Wedge at 0x7f84b44872b0>,
<matplotlib.patches.Wedge at 0x7f84b4486e30>,
<matplotlib.patches.Wedge at 0x7f84b44859c0>],
[Text(-0.7389062157855284, 0.8148727534244291, 'S'),
Text(0.5271737867829742, -0.9654469423686107, 'C'),
Text(1.069316711618934, -0.2579956787476679, 'Q')],
[Text(-0.4030397540648336, 0.4444760473224158, '73.4%'),
Text(0.28754933824525863, -0.5266074231101512, '19.0%'),
Text(0.5832636608830548, -0.1407249156805461, '7.5%')])
```



Fill empty embarked with "S"

```
In [50]: df['Embarked'].fillna('S', inplace=True)
```

```
df['EmbarkedValue'] = df.apply(lambda x: ['S', 'Q', 'C'].index(x['Embarked']), axis=1)
```

Fill "Fare" with mean

```
In [51]: df['Fare'].fillna(value=df['Fare'].mean(), inplace=True)
df['SibSp'].fillna(value=0, inplace=True)
df['Parch'].fillna(value=0, inplace=True)
df['Relatives'] = df['SibSp'] + df['Parch']
df = df.drop(['Parch', 'SibSp'], axis=1)
```

Since the Ticket attribute has 681 unique tickets, it is fair to drop that column

```
In [52]: df['Ticket'].describe()
```

```
Out[52]: count      838
unique      647
top         1601
freq         7
Name: Ticket, dtype: object
```

```
In [53]: df = df.drop(['Ticket'], axis=1)
```


Change column types

```
In [54]: for row in df.columns:
        if df[row].dtype == bool:
            df[row] = df[row].astype(int)
```

Plot correlation matrix

```
In [55]: fig, axis = plt.subplots(figsize=(8, 6))
sns.heatmap(df.corr(numeric_only=True), ax=axis, annot=True)
```

```
Out[55]: <AxesSubplot: >
```




Plot all correlations

```
In [56]: fig, axis = plt.subplots(figsize=(15, 15))
pd.plotting.scatter_matrix(df, ax=axis)
```

To output multiple subplots, the figure containing the passed axes is being cleared.


```
Out[56]: array([[<AxesSubplot: xlabel='Survived', ylabel='Survived'>,
<AxesSubplot: xlabel='Pclass', ylabel='Survived'>,
<AxesSubplot: xlabel='Age', ylabel='Survived'>,
<AxesSubplot: xlabel='Fare', ylabel='Survived'>,
<AxesSubplot: xlabel='Gender', ylabel='Survived'>,
<AxesSubplot: xlabel='EmbarkedValue', ylabel='Survived'>,
<AxesSubplot: xlabel='Relatives', ylabel='Survived'>],
[<AxesSubplot: xlabel='Survived', ylabel='Pclass'>,
<AxesSubplot: xlabel='Pclass', ylabel='Pclass'>,
<AxesSubplot: xlabel='Age', ylabel='Pclass'>,
<AxesSubplot: xlabel='Fare', ylabel='Pclass'>,
<AxesSubplot: xlabel='Gender', ylabel='Pclass'>,
<AxesSubplot: xlabel='EmbarkedValue', ylabel='Pclass'>,
<AxesSubplot: xlabel='Relatives', ylabel='Pclass'>],
[<AxesSubplot: xlabel='Survived', ylabel='Age'>,
<AxesSubplot: xlabel='Pclass', ylabel='Age'>,
<AxesSubplot: xlabel='Age', ylabel='Age'>,
<AxesSubplot: xlabel='Fare', ylabel='Age'>,
<AxesSubplot: xlabel='Gender', ylabel='Age'>,
<AxesSubplot: xlabel='EmbarkedValue', ylabel='Age'>,
<AxesSubplot: xlabel='Relatives', ylabel='Age'>],
[<AxesSubplot: xlabel='Survived', ylabel='Fare'>,
<AxesSubplot: xlabel='Pclass', ylabel='Fare'>,
<AxesSubplot: xlabel='Age', ylabel='Fare'>,
<AxesSubplot: xlabel='Fare', ylabel='Fare'>,
<AxesSubplot: xlabel='Gender', ylabel='Fare'>,
<AxesSubplot: xlabel='EmbarkedValue', ylabel='Fare'>,
<AxesSubplot: xlabel='Relatives', ylabel='Fare'>],
[<AxesSubplot: xlabel='Survived', ylabel='Gender'>,
<AxesSubplot: xlabel='Pclass', ylabel='Gender'>,
<AxesSubplot: xlabel='Age', ylabel='Gender'>,
<AxesSubplot: xlabel='Fare', ylabel='Gender'>,
<AxesSubplot: xlabel='Gender', ylabel='Gender'>,
<AxesSubplot: xlabel='EmbarkedValue', ylabel='Gender'>,
<AxesSubplot: xlabel='Relatives', ylabel='Gender'>],
[<AxesSubplot: xlabel='Survived', ylabel='EmbarkedValue'>,
<AxesSubplot: xlabel='Pclass', ylabel='EmbarkedValue'>,
<AxesSubplot: xlabel='Age', ylabel='EmbarkedValue'>,
<AxesSubplot: xlabel='Fare', ylabel='EmbarkedValue'>,
<AxesSubplot: xlabel='Gender', ylabel='EmbarkedValue'>,
<AxesSubplot: xlabel='EmbarkedValue', ylabel='EmbarkedValue'>,
<AxesSubplot: xlabel='Relatives', ylabel='EmbarkedValue'>],
[<AxesSubplot: xlabel='Survived', ylabel='Relatives'>,
<AxesSubplot: xlabel='Pclass', ylabel='Relatives'>,
<AxesSubplot: xlabel='Age', ylabel='Relatives'>,
<AxesSubplot: xlabel='Fare', ylabel='Relatives'>,
<AxesSubplot: xlabel='Gender', ylabel='Relatives'>,
<AxesSubplot: xlabel='EmbarkedValue', ylabel='Relatives'>,
<AxesSubplot: xlabel='Relatives', ylabel='Relatives'>]],
dtype=object)
```



Plot male-female Age statistic

```
In [57]: fig = sns.FacetGrid(df, hue='Gender', aspect=4)
fig.map(sns.kdeplot, 'Age', fill=True)
oldest = df['Age'].max()
fig.set(xlim=(0, oldest))
fig.add_legend()
```


```
Out[57]: <seaborn.axisgrid.FacetGrid at 0x7f849d833130>
```



Plot male-female Count statistic


```
In [58]: fig, axis = plt.subplots(figsize=(8, 6))
male_female_counts = df['Sex'].value_counts()
axis.pie(male_female_counts, labels=male_female_counts.index, autopct='%1.1f%%')
```

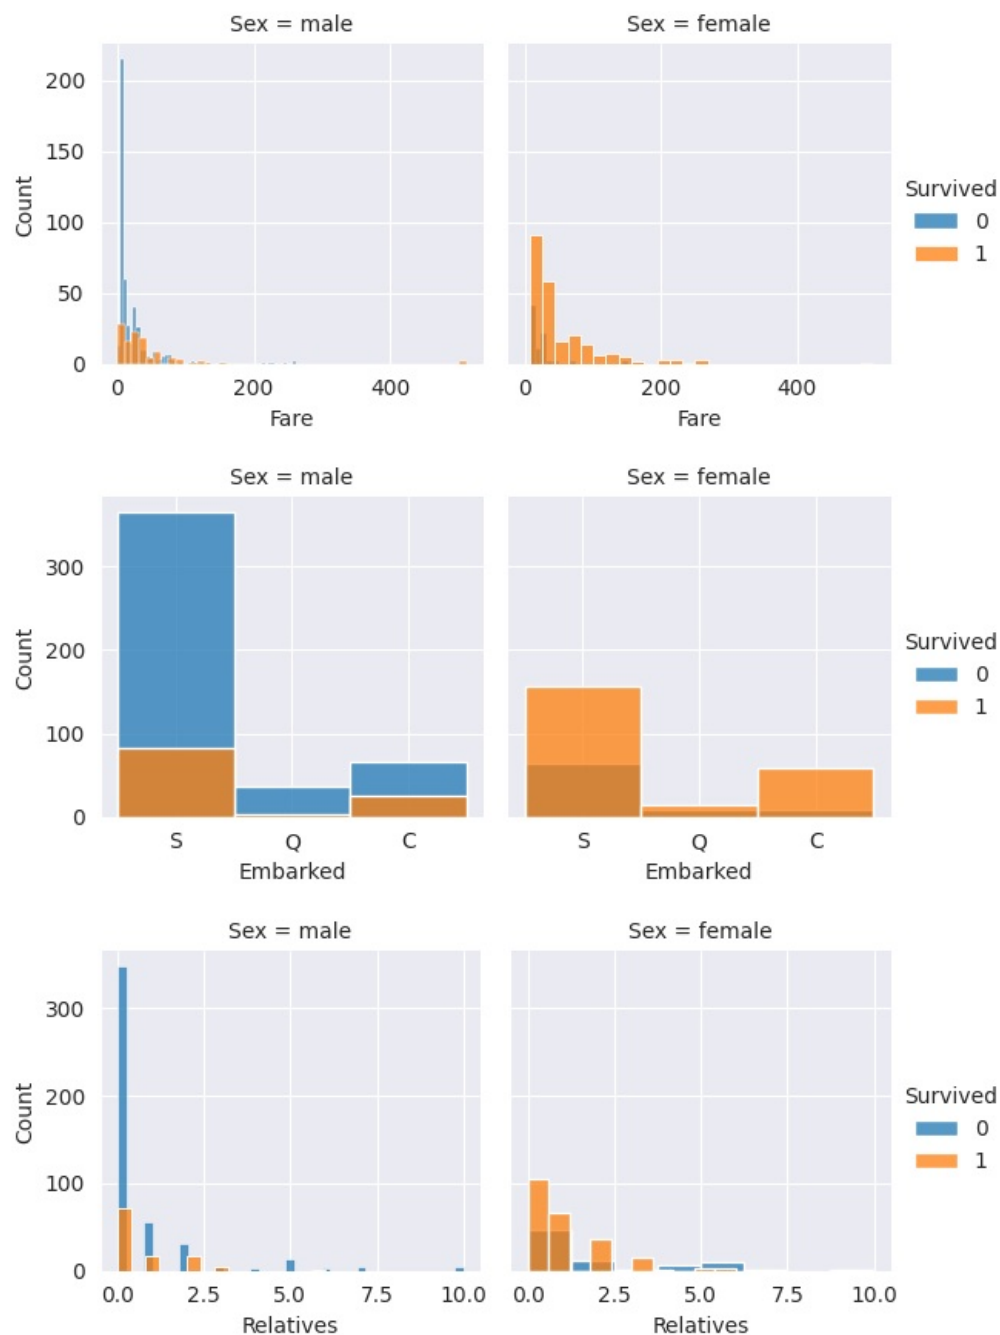
```
factors = df.columns.values.tolist()
factors = [f for f in factors if f not in ['Sex', 'Survived', 'Gender', 'Name', 'EmbarkedValue']]
```



Plot male-female Survival Rate based on Age statistic

```
In [59]: for row, factor in enumerate(factors):
g = sns.FacetGrid(df[['Sex', 'Survived'] + [factor]], col='Sex', hue='Survived')
g.map(sns.histplot, factor)
g.add_legend()
```






KMeans without PCA

```
In [60]: factors.remove('Embarked')
factors.append('EmbarkedValue')
kmeans_kwargs = {
    'init': 'random',
    'n_init': 10,
    'max_iter': 300,
    'random_state': 0,
}
sse = []
max_kernels = 10
features = df[['Survived'] + factors]
for k in range(1, max_kernels + 1):
    kmeans = KMeans(n_clusters=k, **kmeans_kwargs)
    kmeans.fit(features)
    sse.append(kmeans.inertia_)

plt.figure(figsize=(10, 8))
plt.plot(range(1, max_kernels + 1), sse)
```

```
plt.xticks(range(1, max_kernels + 1))
plt.xlabel('Number of Clusters')
plt.ylabel('SSE')
plt.grid(linestyle='--')
```



Elbow point

```
In [61]: kl = KneeLocator(range(1, max_kernels + 1), sse, curve='convex', direction='decreasing')
kl.elbow
```

Out[61]: 3

Set number of clusters to elbow point

```
In [62]: kmeans = KMeans(
    init='random',
    n_clusters=kl.elbow,
    n_init=10,
    max_iter=300,
    random_state=0
)
kmeans.fit(features)
```

```
Out[62]: KMeans
KMeans(init='random', n_clusters=3, n_init=10, random_state=0)
```

```
In [63]: kmeans.cluster_centers_
```

```
Out[63]: array([[ 0.67391304,  1.23913043, 34.96908602, 83.96026449,
  1.41304348,  0.72463768],
 [ 0.32152589,  2.54359673, 29.03663864, 16.70048021,
  0.72752044,  0.35286104],
 [ 0.68421053,  1.        , 31.66986452, 280.19978421,
  1.68421053,  1.15789474]])
```

```
In [64]: fig = px.scatter_3d(
    df, x='Survived', y='Fare', z='Relatives',
    color=kmeans.labels_, hover_data=['Sex', 'Age', 'EmbarkedValue'],
```

```
width=1000, height=800,  
title='Survived-Fare-Relatives KMeans Plot Clusters'  
)  
fig.update(layout_coloraxis_showscale=False)  
fig.show()
```

Agglomerate Clustering

```
In [65]: model = AgglomerativeClustering(n_clusters=kl.elbow, affinity='euclidean', linkage='ward')  
clust_labels = model.fit_predict(features)  
agglomerative = pd.DataFrame(clust_labels)  
agglomerative
```

```
/usr/local/lib/python3.10/dist-packages/sklearn/cluster/_agglomerative.py:983: FutureWarning:  
Attribute `affinity` was deprecated in version 1.2 and will be removed in 1.4. Use `metric` instead
```

Out[65]:

	0
0	1
1	2
2	1
3	1
4	1
...	...
886	1
887	1
888	1
889	1
890	1

891 rows x 1 columns

```
In [66]: df.insert((df.shape[1]), 'agglomerative', agglomerative)
fig = px.scatter_3d(
    df, x='Survived', y='Fare', z='Relatives',
    color=df['agglomerative'], hover_data=['Sex', 'Age', 'EmbarkedValue'],
    width=1000, height=800,
    title='Survived-Fare-Relatives Agglomerative Plot Clusters'
)
fig.update(layout_coloraxis_showscale=False)
fig.show()
```

Results

Результати кластеризації показують невелику різницю у двох методах на даному наборі даних. Можливо, що інформації замало, тому не показується різниця

Additional Assignment

Useful Funcs

```
In [67]: def read_dataset(path: str, sep: str = ';') -> pd.DataFrame:
        data = pd.read_csv(path, sep=sep)
        return data

def replace_comma_with_dots(dataset: pd.DataFrame, column_name: str) -> None:
    dataset[column_name] = dataset[column_name].astype(str)
    dataset[column_name] = dataset[column_name].str.replace(',', '.')

def convert_column_to_float(dataset: pd.DataFrame, column_name: str) -> None:
    dataset[column_name] = dataset[column_name].astype(float)

def replace_nan_with_mean(dataset: pd.DataFrame, column_name: str):
    mean_value = dataset[column_name].mean()
    dataset[column_name].fillna(value=mean_value, inplace=True)

def convert_float_with_positive(dataset: pd.DataFrame, column_name: str):
    dataset[column_name] = dataset[column_name].abs()
```

Read Dataset and filter it

```
In [68]: data_path = 'data/Data2.csv'
df = read_dataset(data_path, sep=';')

df['Population'] = df['Populatiion']
df = df.drop(['Populatiion'], axis=1)

numeric_cols = df.columns[2:]

for column_name in numeric_cols:
    replace_comma_with_dots(df, column_name)
    convert_column_to_float(df, column_name)
    replace_nan_with_mean(df, column_name)
    convert_float_with_positive(df, column_name)
    print(column_name, df[column_name].dtype)

## density column
df['GDP'] = df['GDP per capita'] * df['Population']

df['Density'] = df['Population'] / df['Area']
df.head(10)
```

```
GDP per capita float64
CO2 emission float64
Area float64
Population float64
```

Out [68]:

	Country Name	Region	GDP per capita	CO2 emission	Area	Population	GDP	Density
0	Afghanistan	South Asia	561.778746	9809.225000	652860.0	34656032.0	1.946902e+10	53.083405
1	Albania	Europe & Central Asia	4124.982390	5716.853000	28750.0	2876101.0	1.186387e+10	100.038296
2	Algeria	Middle East & North Africa	3916.881571	145400.217000	2381740.0	40606052.0	1.590491e+11	17.048902
3	American Samoa	East Asia & Pacific	11834.745230	165114.116337	200.0	55599.0	6.580000e+08	277.995000
4	Andorra	Europe & Central Asia	36988.622030	462.042000	470.0	77281.0	2.858518e+09	164.427660
5	Angola	Sub-Saharan Africa	3308.700233	34763.160000	1246700.0	28813463.0	9.533511e+10	23.111786
6	Antigua and Barbuda	Latin America & Caribbean	14462.176280	531.715000	440.0	100963.0	1.460145e+09	229.461364
7	Argentina	Latin America & Caribbean	12440.320980	204024.546000	2780400.0	43847430.0	5.454761e+11	15.770188
8	Armenia	Europe & Central Asia	3614.688357	5529.836000	29740.0	2924816.0	1.057230e+10	98.346200
9	Aruba	Latin America & Caribbean	13374.833168	872.746000	180.0	104822.0	1.401977e+09	582.344444

Factors


In [69]:


```
factors = df.columns.values.tolist()
factors = [f for f in factors if f not in ['Region', 'Country Name']]
```

Frequency diagrams

In [70]:

```
for column_name in factors:
    fig = sns.FacetGrid(df, aspect=4)
    fig.map(sns.kdeplot, column_name, fill=True)
    oldest = df[column_name].max()
    fig.set(xlim=(0, oldest))
    fig.add_legend()
```






Check for linear correlation

```
In [71]: df_nums = df[factors]
fig, axis = plt.subplots(figsize=(8, 6))
sns.heatmap(df_nums.corr(numeric_only=True), ax=axis, annot=True)
```

```
Out[71]: <AxesSubplot: >
```



Linear Func

```
In [72]: def is_linear(series_one: pd.Series, series_two: pd.Series) -> bool:
         return math.fabs(series_one.corr(series_two)) >= 0.8
```

```
rows = []
for c_one in factors:
    corrs_row = [{c: is_linear(df[c_one], df[c])} for c in factors]
    pairs = {}
    for cc in corrs_row:
        pairs.update(cc)
    row = {c_one: pairs}
    rows.append(row)
rows
```




```
Out[72]: [{'GDP per capita': {'GDP per capita': True,
    'CO2 emission': False,
    'Area': False,
    'Population': False,
    'GDP': False,
    'Density': False}},
 {'CO2 emission': {'GDP per capita': False,
    'CO2 emission': True,
    'Area': False,
    'Population': True,
    'GDP': True,
    'Density': False}},
 {'Area': {'GDP per capita': False,
    'CO2 emission': False,
    'Area': True,
    'Population': False,
    'GDP': False,
    'Density': False}},
 {'Population': {'GDP per capita': False,
    'CO2 emission': True,
    'Area': False,
    'Population': True,
    'GDP': False,
    'Density': False}},
 {'GDP': {'GDP per capita': False,
    'CO2 emission': True,
    'Area': False,
    'Population': False,
    'GDP': True,
    'Density': False}},
 {'Density': {'GDP per capita': False,
    'CO2 emission': False,
    'Area': False,
    'Population': False,
    'GDP': False,
    'Density': True}}]
```

Clusterization

```
In [73]: kmeans_kwargs = {
    'init': 'random',
    'n_init': 10,
    'max_iter': 300,
    'random_state': 0,
}
sse = []
max_kernels = 10
features = df[factors]
for k in range(1, max_kernels + 1):
    kmeans = KMeans(n_clusters=k, **kmeans_kwargs)
    kmeans.fit(features)
    sse.append(kmeans.inertia_)

plt.figure(figsize=(10, 8))
plt.plot(range(1, max_kernels + 1), sse)
plt.xticks(range(1, max_kernels + 1))
plt.xlabel('Number of Clusters')
plt.ylabel('SSE')
plt.grid(linestyle='--')
```



Elbow Point

```
In [74]: kl = KneeLocator(range(1, max_kernels + 1), sse, curve='convex', direction='decreasing')
kl.elbow
```

```
Out[74]: 2
```

Plot clusterizations with different number of clusters and region statistic in them

```
In [75]: n_clusters_list = [kl.elbow, 3, 4, 5]
for n in n_clusters_list:
    kmeans = KMeans(
        init='random',
        n_clusters=n,
        n_init=10,
        max_iter=300,
        random_state=0
    )
    kmeans.fit(features)
    name = f'N-Clusters {n}'
    df[name] = kmeans.labels_

    print(name)
    for i in range(n):
        t = df[df[name] == i]
        table = pd.pivot_table(
            t, values=['GDP', 'Population', 'Area'], index=['Region'], aggfunc=np.sum
        )
        table['Density'] = table['Population'] / table['Area']
        table['GDP per capita'] = table['GDP'] / table['Population']
        max_gdp = table['GDP per capita'].max()
        max_density = table['Density'].max()
        max_gdp_region = table[table['GDP per capita'] == max_gdp].index[0]
        max_density_region = table[table['Density'] == max_density].index[0]
        print(f'\tCluster {i}')
        print(f'\t\tMax GDP: {max_gdp}\tRegion: {max_gdp_region}')
        print(f'\t\tMax Density: {max_density}\tRegion: {max_density_region}')
```

```

fig = px.scatter_3d(
    df, x='GDP per capita', y='Area', z='Population',
    color=kmeans.labels_, hover_data=['Country Name', 'Region', 'CO2 emission', 'Density'],
    width=1000, height=800,
    title=f'N-Clusters {n}'
)
fig.update(layout_coloraxis_showscale=False)
fig.show()

```

N-Clusters 2

```

Cluster 0
    Max GDP: 42131.52083010065    Region: North America
    Max Density: 343.9666814206596    Region: South Asia
Cluster 1
    Max GDP: 57638.15909    Region: North America
    Max Density: 144.1679212532669    Region: East Asia & Pacific

```

N-Clusters 3

```

Cluster 0
    Max GDP: 16366.379137353946    Region: Europe & Central Asia
    Max Density: 1306.62    Region: North America
Cluster 1
    Max GDP: 42183.2951    Region: North America
    Max Density: 402.8192953460619    Region: South Asia
Cluster 2
    Max GDP: 57638.15909    Region: North America
    Max Density: 144.1679212532669    Region: East Asia & Pacific

```

N-Clusters 4

Cluster 0

Max GDP: 14626.78085289752 Region: Europe & Central Asia
Max Density: 1306.62 Region: North America

Cluster 1

Max GDP: 42183.2951 Region: North America
Max Density: 34.50302001624005 Region: East Asia & Pacific

Cluster 2

Max GDP: 39965.22390624113 Region: Europe & Central Asia
Max Density: 402.8192953460619 Region: South Asia

Cluster 3

Max GDP: 57638.15909 Region: North America
Max Density: 144.1679212532669 Region: East Asia & Pacific

N-Clusters 5

Cluster 0

Max GDP: 37771.51075860539 Region: Europe & Central Asia
Max Density: 377.3908172888432 Region: South Asia

Cluster 1

Max GDP: 42183.2951 Region: North America
Max Density: 34.50302001624005 Region: East Asia & Pacific

Cluster 2

Max GDP: 39965.22390624113 Region: Europe & Central Asia
Max Density: 402.8192953460619 Region: South Asia

Cluster 3

Max GDP: 57638.15909 Region: North America
Max Density: 144.1679212532669 Region: East Asia & Pacific

Cluster 4

Max GDP: 13374.83316831 Region: North America
Max Density: 1306.62 Region: North America

Results

Проаналізували корельованість параметрів, статистику решіонів в кластерах та намалювали частотні діаграми